

Dhaga is acquiring customers faster than it can recover value from orders.

A bounded MVP turns unstructured return feedback into actionable, human-reviewed decisions.

- 31% overall return rate creates substantial operational pressure.
- 44% of return reasons are “Other”, limiting prioritization.
- Financial recovery is a pilot hypothesis—not a proven result.

Central thesis · diagnose first, intervene safely, measure verified recovery.



31% returns is the burden; "Other" is the diagnosis gap.



48,000

weekly orders

Case-study
baseline.

31%

overall return rate

Case-study
baseline.

14,880

implied weekly
returns

Calculated:
 $48,000 \times 31\%$.

44% of returns are "Other" → approximately 6,547 comments/week (calculated: $14,880 \times 44\%$). The gap is a starting point for diagnosis, not proof that all returns are preventable.

“Other” is a diagnosis gap, not a root cause.

The business decision is not “how do we automate returns?” It is “which return problems are worth fixing first?”

Repeat purchase 22% has been flat for six quarters: a case-study signal, not a causal claim.	Acquisition economics CAC is up 40% YoY, increasing pressure to recover value from each order.
Reverse logistics COD RTO is 26% and reverse logistics is about ₹120 per case in the baseline context.	GMV context ₹310 Cr GMV frames the opportunity size; it does not prove recoverable value.

Different causes imply different interventions. The highest-volume reason is not necessarily the highest-cost reason; prioritize using both frequency and financial impact.

What happens when we understand the reason?

Return signal	What we learn	Possible intervention
Fit / Size	Whether the product or guidance missed the customer’s fit need.	Size-chart/catalog fix or eligible exchange.
Colour mismatch	Whether imagery or shade description set the wrong expectation.	Imagery or shade-description correction.
Listing mismatch	Whether product content failed to match the delivered item.	Listing/content correction.
Quality issue	Whether the issue points to product, vendor, or category quality.	Vendor/category investigation.
Non-product	Whether the issue belongs with operations, logistics, or CX.	Route to the appropriate operational team.

The MVP doesn’t solve every problem. It makes the right problem visible to the right owner. The highest-volume reason is not necessarily the highest-cost reason. Prioritize causes using both frequency and financial impact.

What does the MVP actually change?

BEFORE

Other

“Size M fit nahi aa raha, exchange ke liye size L chahiye.”

Manual interpretation, inconsistent routing, no structured reason.

AFTER

Fit/Size + confidence + suggested action

Category/CX reviewer approves, edits, or dismisses.

The AI doesn't make the decision. It prepares the decision.

The 75% confidence threshold remains configurable prototype routing, calibrated against manually labelled data.

Start with return classification because it connects diagnosis to a decision.

01 · PRIORITIZE

Return classification

Directly addresses the 44% “Other” gap, creates a usable reason signal, and keeps review human.

02

Support triage

Useful, but closer to a symptom than the underlying return-value leak.

03

Listing acceleration

Potentially valuable, but depends on knowing which mismatches matter.

04

Vendor PO visibility

Important foundation work, but a slower path to the immediate diagnosis problem.

The MVP is the smallest useful intervention: classification, confidence, suggested action, reviewer control, and outcome capture. No unsupported scoring or causal claim.

MVP SCOPE · BASELINE → INTERVENTION

Classify the reason; keep every consequential action human.

1

Ingest + validate

Hinglish/free-text comments; validate required inputs and context.

2

Classify + suggest

Map to candidate categories; generate confidence and a suggested action.

3

Route + learn

Route ambiguous/low-confidence cases; capture corrections and outcomes.

75% confidence is a configurable prototype routing rule, not a validated accuracy guarantee. Calibrate it against manually labelled data.

AI recommends. Category/CX decides. No automatic customer messaging, exchange execution, refund, or return decision.



From “Other” text to a reviewable decision queue.

A thin path connects the existing return signal to structured insight, human review, and measurable outcomes.

01 INGEST

Return text

Hinglish/free-text comments and available order context.

02 VALIDATE

Check inputs

Confirm usable text and required context.

03 CLASSIFY

Schema + confidence

Fit/size, colour mismatch, listing mismatch, quality issue, non-product, or other verified class.

04 REVIEW

Approve / edit / dismiss

Category/CX owns the decision; corrections feed learning.

Architecture supports a review queue, not autonomous operations. Downstream exchange, message, refund, or return action stays outside the MVP boundary.

The screen makes the human decision visible.

OVERVIEW

What the prototype displays

- 31% return rate and 44% "Other" context.
- 21 auto-approved classification records with displayed 91% status.
- 2 open review records and 11.2-second displayed agent SLA.

These are prototype display values, not verified business results.

REVIEW QUEUE

Where the operator acts

- Inspect category, confidence, and context.
- Approve, edit, or dismiss the recommendation.
- Keep customer and operational actions behind approval.

"Auto-approved" means classification/prototype status only.

AI recommends. Category/CX decides. Use the latest portal snapshot as prototype context; do not present its displayed metrics as pilot proof.

Where value recovery happens

01

6,547

"Other" comments/week

02

Classify

Identify the likely cause.

03

Review

Human reviews recommendation.

04

Intervene

Listing fix, product/vendor fix, or eligible exchange.

05

Measure

Outcome, then verified cost avoided/value retained.

Classification is the first step, not the outcome. Prototype exchange/freight figures are simulated or modelled hypotheses, never verified business results.

Prototype output is not the same as validated business outcome.

Prototype / model outputs

- 21 auto-approved records; displayed 91% status.
- 2 open review records; 11.2s displayed agent SLA.
- 74% simulated doorstep-exchange conversion.
- ₹42,280 displayed freight saved.
- ₹5.81 lakh/week displayed freight preserved in the funnel.
- ₹40.70 lakh/week displayed GMV retained.
- Earlier view: ₹51,081/week displayed; views require reconciliation.

Pilot outcomes to verify

- Accuracy by category and actionable-category rate.
- Correction/escalation rate and review time per case.
- Offers, acceptances, and completed exchanges separately.
- Verified reverse-logistics cost avoided, contribution margin, and GMV retained with documented definitions.

Auto-approved is classification status only. Prototype/model values are simulated or displayed outputs, not verified business results. Reconcile the financial views before any impact claim.

Measure the chain from diagnosis to verified recovery.

WEEK 1 · BASELINE + CALIBRATION

Know the starting point

- Select one category and a representative sample.
- Manually label reasons; establish return, review-effort, and reverse-cost baselines.
- Calibrate the configurable 75% routing threshold against labelled data and agree success criteria.

WEEK 2 · RUN + REVIEW + EVALUATE

Test the intervention

- Classify, review, and record corrections, decisions, and review time.
- Track exchange offers, acceptances, and completed exchanges separately.
- Calculate verified cost avoided; do not promise a two-week return-rate reduction.

Pilot measures: accuracy by category · actionable-category rate · correction/escalation rate · review time · exchange offers/acceptances/completions · verified reverse-logistics cost avoided · contribution margin and GMV retained with documented definitions.

Approve one category, then scale only on verified evidence.

Approve

One category for a two-week Hinglish/free-text classification pilot.

Sign off baseline fields and nominate the Category/CX owner.

Measure

Accuracy, actionability, corrections, review time, and exchange stages.

Verify reverse cost and margin/GMV definitions before reporting impact.

Scale / stop

Scale only if client-agreed criteria are met without unacceptable exception or refund load.

Stop or refine if data quality, accuracy, or economics fail.

Boundary

Human approval remains between classification and any action.

No automatic messages, exchanges, refunds, or return decisions.